

# CHAPTER 20: MULTITHREADING AND MULTITASKING

In this chapter, we will break down how Windows handles **multitasking** and **multithreading**, with simple explanations, key concepts, and practical examples. We'll also include insights from Charles Petzold's book to give context.

## 1. Multitasking

**Multitasking** is the ability of the operating system to run multiple programs at the same time.

- Windows does this by giving each program a **time slice**.
- Programs appear to run simultaneously, even if the CPU is only executing one instruction at a time.
- This improves **system responsiveness** and lets you switch between apps seamlessly.

### Evolution in Windows:

- **16-bit Windows:** Used **cooperative multitasking**. Programs had to voluntarily give up CPU control. If one program “hung,” the whole system could freeze.
- **32-bit Windows:** Introduced **preemptive multitasking**. The OS actively manages CPU time, making sure no program monopolizes it.

## 2. Multithreading

**Multithreading** is when a single program splits into multiple **threads**, which are lightweight units of execution within the program.

### Benefits of multithreading:

- Run **background tasks** without freezing the interface.
- Keep **UI responsive** while other operations run.
- Execute **independent tasks in parallel**, improving performance on multiprocessor systems.

### 3. Key Terminology

- **Process:** A running program with its own memory and resources.
- **Thread:** A smaller execution unit inside a process. Threads share memory and resources with their parent process.
- **Context Switching:** Saving and restoring a thread's state when switching to another thread.
- **Synchronization:** Methods to safely coordinate access to shared resources between threads, avoiding **race conditions** or data corruption.

### 4. Topics Covered in Chapter 20

#### Thread Creation and Management

- Use CreateThread to start a new thread.
- Threads can have **priorities** set, and you can **suspend**, **resume**, or **terminate** them.

#### Synchronization Techniques

- Prevent conflicts when threads access shared data using:
  - ✓ **Critical Sections** – fast, process-local locks.
  - ✓ **Mutexes** – system-wide locks.
  - ✓ **Semaphores** – control access to a resource with limited availability.
  - ✓ **Events** – signal between threads.

#### Thread-Specific Storage

- Use TlsAlloc, TlsGetValue, and TlsSetValue to give each thread its own private data.
- Useful when multiple threads need independent copies of the same type of data.

#### Win32 Timers

- SetTimer and KillTimer schedule recurring or one-time events.
- Timers are handy for repeating tasks without blocking the main thread.

## Asynchronous Procedure Calls (APC)

- Execute code asynchronously in a separate thread using:
  - ✓ BeginThreadEx – create a thread safely in Windows.
  - ✓ QueueUserAPC – queue code to run in a specific thread.

## 5. Best Practices in Multithreading

- Avoid **deadlocks** by careful locking order.
- Optimize **thread performance** by limiting unnecessary threads.
- Ensure **thread safety** when multiple threads access shared resources.

## MULTITASKING IN THE DOS ERA

Before Windows made multitasking smooth, DOS had **ideas** but a lot of **practical limits**.

- Users **wanted multitasking** (running multiple programs at once).
- DOS **wasn't built for it** — hardware and software got in the way.
- Early attempts were **creative but limited**.

## Why DOS Struggled

### I. Hardware Limits


- CPUs like the **Intel 8088** had **no memory management unit (MMU)**.
- Programs had to share **one flat memory space**, making juggling apps tricky.
- Moving memory blocks to free up space was slow and hard.

### II. DOS Architecture

- DOS was **simple and lean** — designed for **single-tasking**.
- APIs were basic: mostly **file access** and **program loading**.
- No robust system services → developers couldn't easily implement multitasking.

Think of DOS as a **tiny studio apartment**:

- Only one person (program) can comfortably live there at a time.
- You *could* squeeze in another, but furniture (memory) had to be constantly moved.
- Windows later gave everyone **their own rooms (protected memory & system services)** — suddenly multitasking made sense.



# Creative Workarounds & Multitasking in DOS 🖥️🌟

Even though DOS wasn't designed for multitasking, **clever programmers found ways to simulate it.**

## 1. TSR Programs (Terminate-and-Stay-Resident)

- TSRs stayed **in memory** after their main program closed, running in the background.
- Examples:
  - ✓ **Print spoolers:** used the hardware timer interrupt to print while you worked on another program.
  - ✓ **Borland SideKick:** temporarily paused your current app to show its interface, then returned control.
- TSRs were **early multitasking hacks**, letting small tools coexist with other programs.

## 2. Enhanced DOS Features

- Microsoft added **memory swapping to disk** and other tweaks.
- These **helped manage memory** better, indirectly supporting background activities.

## 3. Market Experiments

- **Task-switching shells** like **Quarterdeck DesqView** allowed running multiple DOS programs.
- Reality check: performance was slow, setup was tricky, adoption was limited.

## 4. Key Takeaways

- Multitasking was **desired by users**, even on single-user PCs.
- Programmers used **TSRs and simple task-switchers** to stretch DOS's limits.
- These efforts **highlighted the need for OS-level multitasking**, eventually realized in Windows.

## 5. Exploration Ideas

- Study **popular TSRs**: what they did, how they worked.
- Dive into **DOS memory management** and **context switching** challenges.
- Compare DOS multitasking hacks vs. **Windows 3.x and later**, seeing how true multitasking became practical.

💡 Mental shortcut: DOS multitasking = **hacks in a single-room apartment**.

- TSRs = roommates who quietly exist in the background.
- Task-switchers = rotating chairs; clunky but functional.
- Windows = actual multiple rooms, each app gets space, fully coordinated.

## Multitasking in Early Windows ⚡

Early Windows (1.0, 1985) **brought multitasking out of DOS limitations**, giving multiple programs a graphical workspace and basic coordination.

### 1. Windows 1.0's Multitasking Breakthrough

- **Graphical interface**: Multiple programs could run concurrently, unlike DOS command-line shells.
- **Cooperative multitasking (nonpreemptive)**:
  - ✓ Programs only yield control voluntarily after processing messages.
  - ✓ Relies on well-behaved programs; a stuck program could freeze the system.
- **Message-based architecture**: Programs are idle until a message arrives.

### Workarounds & Limitations

- **Preemption**: Used only for DOS programs or certain multimedia tasks (DLLs needing hardware timing).
- **Hourglass cursor**: Visual cue for busy programs.
- **Windows timer & PeekMessage**: Let programs periodically handle messages, preventing total freeze during long tasks.

## 2. Data Sharing Mechanisms

### a) Clipboard

- **Basic & versatile:** For cut/copy/paste across applications.
- **Temporary storage:** Holds data for short-term transfer.
- **Manual operation:** User-controlled, not automatic.

### b) Dynamic Data Exchange (DDE)

- **Live links:** Programs exchange data in real-time, even if idle.
- **Client-server model:** One app requests updates from another.
- **Examples:** Stock tickers, spreadsheets linked to databases.
- **Caveats:** Complex, fragile if connections fail.

### c) Object Linking and Embedding (OLE)

- **Embedding objects:** One document can contain content from another program.
- **In-place editing:** Edit objects directly within the main document.
- **Example:** Edit a spreadsheet chart inside Word without launching Excel.
- **Use case:** Rich, interactive compound documents.

## 3. Key Points

- Clipboard → simple, manual data sharing.
- DDE → live updates between programs, more complex.
- OLE → seamless object integration for richer documents.
- 16-bit Windows used **cooperative multitasking**, with no enforced preemption.
- System responsiveness relied on programs yielding control voluntarily.

## 4. Further Exploration

- Investigate challenges of **preemptive multitasking** in 16-bit Windows.
- Study real examples of **application freezes or bottlenecks** under cooperative multitasking.
- Track the evolution toward **preemptive multitasking** in 32-bit Windows.
- Modern parallels: **Clipboard history, cloud sync, XML/JSON** for universal data exchange, and **OLE alternatives** like ActiveX or .NET components.

### 💡 Mental shortcut:

- DOS multitasking = roommates squeezing in one apartment.
- Windows 1.0 = a shared living room where everyone must **take turns politely**.
- Clipboard/DDE/OLE = ways to **pass notes between roommates without chaos**.

## Multithreading & The Evolution of Input

### 1. The History Lesson: Why OS/2 Failed

Before Windows 95/NT, Microsoft and IBM built **OS/2 Presentation Manager (PM)**. It was a 32-bit operating system with a fatal flaw in how it handled user input.

**The Flaw: The Serialized Message Queue** In OS/2, the system had *one* single pipe for all keyboard and mouse clicks for every running program.

- **The Rule:** The system would not process Input B until the application finished processing Input A.
- **The Goal:** Predictability. If you type "ABC", the system guarantees "A" is processed before "B".
- **The Reality:** If one program crashed or hung while processing "A", the entire system froze. You couldn't click on another window because the mouse click was stuck in the queue behind "A".

**The Lesson:** A robust OS cannot let one bad app freeze the whole mouse.

## 2. The Windows Solution: Deserialized Input

Modern Windows (Win32 API) fixed this by giving every thread its own private message queue.

- **How it works:** When you click on Chrome, Windows looks at where the mouse is, determines which thread owns that window, and drops the message directly into *that specific thread's* queue.
- **The Benefit:** If Chrome hangs, Notepad keeps working. You can Alt-Tab away from a frozen program.

## 3. The Architecture of a Multithreaded App

In modern Windows programming, we divide labor to keep the application responsive.

### The "Governor & Staff" Model

#### 1. The Primary Thread (The Governor):

- ✓ **Job:** Creates windows, runs the Message Loop, handles the UI (WM\_PAINT, buttons).
- ✓ **Rule:** Never do heavy lifting here. If you calculate Pi to the billionth digit here, the UI freezes, and the window says "(Not Responding)".

#### 2. The Secondary Threads (The Staff):

- ✓ **Job:** Long calculations, file I/O, networking.
- ✓ **Rule:** They do *not* own windows. They crunch numbers in the background and tell the Governor when they are done.

## 4. What Threads Share vs. What They Keep

When you create a new thread, it is not a separate program. It lives inside the same "house" (Process).

THREAD MEMORY MODEL: SHARED VS. PRIVATE	
 THE SHARED HOUSE (Process Scope)	 THE PRIVATE BACKPACK (Thread Scope)
<b>Global Variables</b> All threads see and can modify global data. Requires synchronization (Mutex/CriticalSection).	<b>The Stack</b> Each thread has its own function call history, return addresses, and local <code>auto</code> variables.
<b>Heap Memory</b> Pointers from <code>malloc</code> or <code>new</code> are accessible by everyone if the address is shared.	<b>CPU Registers</b> Context includes the Instruction Pointer (EIP/RIP) and Stack Pointer (ESP/RSP).
<b>System Resources</b> Open files, kernel handles, and Window objects (HWND) belong to the process.	<b>Thread Local Storage (TLS)</b> Special <code>__declspec(thread)</code> variables that are unique instances for every thread.

**The Danger:** Because they share global memory, two threads can try to write to the same variable at the exact same time, causing corruption (Race Conditions).

---

## 5. Summary Checklist

- **OS/2 PM** failed because one frozen app froze the whole system (Serialized Input).
- **Windows** succeeds because input is split per thread (Deserialized Input).
- **Primary Thread:** Handles UI/Messages.
- **Secondary Thread:** Handles heavy math/background work.
- **Threads share memory**, which makes them fast but dangerous (requires synchronization).


**Next Step:** Now that we know *why* we need threads, we need to look at **how to create them** using `_beginthreadex` vs `CreateThread`.

# THREADING CHALLENGES & BEST PRACTICES

Multithreading is powerful, but it introduces a new class of bugs that are arguably the hardest to debug in all of computer science.

## 1. The Chaos: Race Conditions and Deadlocks

In a single-threaded program, you know exactly which line of code runs next. In a multithreaded program, the Operating System acts like a chaotic referee. It pauses Thread A and runs Thread B at random times—sometimes right in the middle of a calculation.



**The Race Condition** This happens when two threads try to change the same variable at the same time.

- *Thread A* reads Score = 10.
- *Thread B* reads Score = 10.
- *Thread A* adds 1 and writes 11.
- *Thread B* adds 1 and writes 11.
- *Result:* The score should be 12, but it is 11. Data is lost because they "raced" to write.

**The Deadlock** This happens when two threads stop and wait for each other.

- *Thread A* holds Key 1 and waits for Key 2.
- *Thread B* holds Key 2 and waits for Key 1.
- *Result:* They both wait forever. The program freezes.

## 2. The Fix: Synchronization

To stop the chaos, you need **Synchronization Primitives**. These are tools that force threads to wait their turn.

- **Critical Sections:** A block of code that only one thread can enter at a time. It is like a bathroom with a lock. If Thread A is inside, Thread B must wait outside until A leaves.
- **Semaphores:** Like a bouncer at a club. It allows a specific number of threads in at once (e.g., allowing 5 connections to a database).

## 3. The Hardware Reality: 16-bit vs 64-bit

**The Old Days (16-bit)** Processors were simple. Even adding 1 to a large number required two CPU cycles. If the OS interrupted the thread between cycle 1 and 2, the number would be corrupted. You had to lock everything.

### The Modern Era (64-bit)

- **Atomicity:** Modern CPUs can read and write huge numbers (64-bit) in a single cycle. You don't need locks just to read a variable.
- **The New Danger (Optimization):** Modern CPUs are smart. They reorder your instructions to run faster. They might run Line 10 before Line 5 if they think it is safe. In multithreading, this can break your logic.
- **Takeaway:** Do not rely on hardware tricks. Always use proper Synchronization (locks) to be safe.

## 4. Windows Advancements

Windows evolved to make threading easier and safer.

**A. Deserialized Input (The Fix for Freezing)** - We covered this in the last section. Windows gives every thread its own input queue so a frozen background thread doesn't freeze the mouse cursor.

**B. Thread Local Storage (TLS)** Global variables are dangerous because all threads share them (leading to Race Conditions). TLS allows you to create a "Global" variable that is unique to each thread.

- If Thread A writes "Red" to *Color*, it sees "Red".
- If Thread B writes "Blue" to *Color*, it sees "Blue". They use the same variable name but look at different memory addresses.

## 5. The "New & Improved" Fallacy

Just because you *can* use threads doesn't mean you *should*. Threads add overhead. The CPU wastes time switching context between them.

**The 1/10 Second Rule (100 Milliseconds)** Use this rule of thumb to decide if you need a thread:

- **Task takes < 100ms:** Do not use a thread. Just run it. The user won't notice the tiny pause.
- **Task takes > 100ms:** Use a secondary thread. If you don't, the UI will freeze, and the user will think the app crashed.

## 6. Summary Checklist

1. **Race Conditions** happen when threads share data without locks.
2. **Deadlocks** happen when threads wait for each other.
3. **Critical Sections** prevent these bugs by forcing single-file access.
4. **TLS** gives threads private data.
5. **Only thread** if the task is slow (over 1/10th of a second).

## The Two Ways to Spawn a Thread

### a) The Raw Windows API (`CreateThread`)

This is the native function provided by the OS.

- **Pros:** It gives you granular control (Security attributes, Stack size).
- **Cons:** It does **not** set up the C Runtime (CRT).
- **The Danger:** If you use C functions like malloc, printf, or strtok inside a thread created with CreateThread, the program might crash or leak memory because the CRT data structures weren't initialized for that thread.

## b) The C Runtime Helper (`_beginthreadex`)

This is the wrapper function provided by Microsoft's C library.

- **Pros:** It initializes the C Runtime, then calls CreateThread internally.
- **Cons:** Slightly different syntax.
- **The Rule:** Always use `_beginthread` (or `_beginthreadex`) if your thread uses any C library functions.

### Comparison Table:

THREAD CREATION: WIN32 API VS. C RUNTIME		
FEATURE	CREATETHREAD	<code>_BEGINTHREADEX</code>
<b>Origin</b>	Windows Native API ( <code>windows.h</code> )	C Runtime Library ( <code>process.h</code> )
<b>Use Case</b>	Pure API coding (No C/C++ Libs)	<b>Standard C/C++ coding</b>
<b>CRT Safety</b>	<b>Unsafe:</b> May cause leaks if using <code>malloc</code> , <code>printf</code> , or <code>strtok</code> .	<b>Safe:</b> Correctly initializes per-thread CRT data blocks.
<b>Returns</b>	<code>HANDLE</code>	<code>uintptr_t (cast to HANDLE)</code>
<b>Exit Method</b>	<code>ExitThread()</code>	<code>_endthreadex()</code>

## The Random Rectangles Program (RNDRCTMT.C)

This program demonstrates the simplest possible multithreaded app:

- **Thread 1 (Main):** Handles the Window (Resizing, Closing).
- **Thread 2 (Worker):** Draws random colored rectangles on the window background forever.

### I. How it works:

1. **WinMain:** Registers the class and creates the window.
2. **WM\_CREATE:** The main thread calls `_beginthread(Thread, ...)` to spawn the worker.
3. **The Worker Loop:**
  - It sits in a `while(TRUE)` loop.
  - It generates random x, y, color values.
  - It calls Rectangle to draw on the screen.

4. **WM\_SIZE**: When you resize the window, the Main Thread updates the global variables cxClient and cyClient. The Worker Thread reads these new values instantly and starts drawing in the new area.

## II. The Compiler Setting (Crucial!)

You cannot just compile this code normally. You must tell the compiler: "*I am using threads.*"

- **The Switch:** /MT (Multithreaded Static) or /MD (Multithreaded DLL).
- **The Library:** It links against LIBCMT.LIB.
- **What it does:** It changes standard functions like strtok. In a single-threaded app, strtok uses a static variable to remember where it left off. In a multithreaded app, LIBCMT replaces that static variable with **Thread Local Storage (TLS)** so two threads using strtok at the same time don't corrupt each other's strings.

## III. Critical Bug Warning (Synchronization)

The RNDRCTMT.C example is simple, but it has a **Hidden Race Condition**.

**The Shared Data:** cxClient and cyClient (Window Size).

**The Race:**

- Thread 1 (Main) is writing cxClient = 500.
- Thread 2 (Worker) is reading cxClient.

It is possible for Thread 2 to read the variable *while* Thread 1 is halfway through writing it (on 16-bit systems especially).

**The Fix:** Real programs need a **Critical Section** to lock the variable while reading/writing.

## IV. Summary Checklist

1. Use \_beginthread instead of CreateThread to keep malloc safe.
2. Enable the **Multithreaded (/MT)** setting in your compiler.
3. **Global Variables** are the easiest way for threads to talk, but they are dangerous without locks.
4. **Automatic Variables** (inside functions) are safe; every thread gets its own copy on its own stack.

## The Multitasking Contest (MULTI1 vs MULTI2)

This section compares two ways to solve a classic 1986 programming problem: "How do you run 4 distinct tasks in 4 separate windows at the same time?"

The Tasks:

1. Window 1: Count up (1, 2, 3...).
2. Window 2: Find Prime Numbers.
3. Window 3: Calculate Fibonacci Sequence.
4. Window 4: Draw random circles.




Multi1 and Multi2  
Program.mp4

### 1. The Old Way: MULTI1 (The Simulation)

**The Strategy: The Juggler** MULTI1 does not use threads. It uses a **Timer**. It relies on a single main loop. Every 10 milliseconds, the timer fires a WM\_TIMER message. The program catches this message and quickly updates Window 1, then Window 2, etc.

**Why it works:** Computers are fast. If you switch between tasks quickly enough, it *looks* like they are happening at the same time.

**The Flaw:** This is "Cooperative Multitasking." If Window 2 gets stuck calculating a massive Prime Number, Windows 3 and 4 stop updating. The entire application freezes until the calculation is done.




## 2. The New Way: MULTI2 (True Multithreading)

The Strategy: The Team MULTI2 creates 5 threads total.

1. **Main Thread:** Handles the UI (moving windows, clicking menus).
2. **4 Worker Threads:** Each window gets its own dedicated thread created with `_beginthread`.

### The Differences:

- **No Timer:** The worker threads don't wait for a "tick." They run while loops as fast as the CPU allows.
- **Struct Passing:** Since `_beginthread` only accepts one argument, we pack the data (Window Handle, Window Size, Kill Flag) into a PARAMS structure and pass the pointer.
- **Responsiveness:** If the Prime Number thread gets stuck on a hard calculation, the Circle thread keeps drawing. The Main thread keeps responding to the mouse.



### 3. The Danger: The "bKill" Race Condition


MULTI2 introduces a serious bug that MULTI1 didn't have.

**The Scenario:** You have a boolean flag called bKill. When the user closes the window, the Main Thread sets bKill = TRUE. The Worker Thread checks this flag; if it is TRUE, it stops.

#### The Crash (Race Condition):

1. Worker Thread checks bKill. It is **FALSE**.
2. *Context Switch happens.* The OS pauses the Worker and runs the Main Thread.
3. User closes the window. Main Thread destroys the window handle and sets bKill = **TRUE**.
4. *Context Switch happens.* The OS resumes the Worker.
5. Worker Thread (thinking bKill is still false) tries to draw on the window.
6. **CRASH.** The window handle is invalid.

**The Lesson:** You cannot rely on simple boolean flags to stop threads safely. In a real application, you must use **Synchronization Objects** (like Event Objects or Critical Sections) to ensure the window isn't destroyed while a thread is painting on it.




## 4. GDI and Threads

Windows GDI (Graphics Device Interface) has specific rules for threads:

- **Windows 98/Legacy:** Not thread-safe. If two threads try to draw at the exact same time, the system might crash.
- **Windows NT/Modern:** The OS "Serializes" GDI calls. It forces Thread B to wait if Thread A is drawing. This prevents crashes but slows performance.

**Best Practice:** Do not share GDI Objects (Pens, Brushes, Device Contexts) between threads. Give every thread its own tools.



## Summary Checklist

1. **MULTI1** fakes multitasking using WM\_TIMER. It is safe but can freeze if one task is heavy.
2. **MULTI2** uses \_beginthread for real multitasking. It is fast but dangerous.
3. **The Struct:** Use a structure to pass multiple arguments to a new thread.
4. **The Race:** Checking a variable (bKill) and acting on it is not an "Atomic" operation. The OS can interrupt you in the middle.
5. **The Fix:** Real multithreading requires **Synchronization** to prevent accessing dead memory.

## SLEEP, SYNCHRONIZATION, AND CRITICAL SECTIONS

We established that multithreading causes chaos (Race Conditions). Now, let's look at the tools we use to tame that chaos.

### 1. The Sleep Function: The "Pause" Button

Sleep(milliseconds) is the simplest way to control a thread.

- **What it does:** It tells the Operating System: "*Stop running me for X milliseconds. Let other threads use the CPU.*"
- **The Sleep(0) Trick:** If you call Sleep(0), you tell the OS: "*I don't need to pause, but I am willing to give up the rest of my turn if anyone else is waiting.*" It is a polite way to yield the processor.

#### When to use it:

- **Good:** In a background thread loop to prevent it from eating 100% CPU (e.g., check for new email every 5 seconds).
- **Bad:** In the Main Thread. Calling Sleep here freezes your window. The user cannot click anything until it wakes up.

## 2. The Traffic Jam: Why We Need Locks

Imagine a bank account with \$100.

- **Thread A** tries to withdraw \$10.
- **Thread B** tries to withdraw \$10.

If they run at the exact same time:

1. Thread A reads balance (\$100).
  2. Thread B reads balance (\$100).
  3. Thread A writes new balance (\$90).
  4. Thread B writes new balance (\$90).
- **Result:** You withdrew \$20, but the balance only dropped by \$10. The bank lost money.

We need a way to force Thread B to wait until Thread A is completely finished.

## 3. The Solution: Critical Sections

A Critical Section is a traffic light for code. It protects a specific block of memory.

**How it works (The 4 Steps):**

1. **Initialize:** Create the "Traffic Light" object (CRITICAL\_SECTION).
2. **Enter:** Before you touch the shared data, call EnterCriticalSection.
  - ✓ *Effect:* If the light is Green, you pass. The light turns Red.
  - ✓ *Effect:* If the light is Red (someone else is inside), your thread **sleeps** immediately. You wait until it turns Green.
3. **Leave:** When you are done, call LeaveCriticalSection.
  - ✓ *Effect:* The light turns Green. If another thread was waiting, it wakes up and enters.
4. **Delete:** When the program ends, clean up the object.

The Code Example:

```
CRITICAL_SECTION cs; // The Traffic Light

// ... Initialization somewhere ...

void UpdateBankAccount() {
    EnterCriticalSection(&cs); // STOP! Wait your turn.

    // --- SAFE ZONE ---
    // Only one thread can be here at a time.
    int balance = ReadBalance();
    balance = balance - 10;
    WriteBalance(balance);
    // -----

    LeaveCriticalSection(&cs); // Go ahead, next person.
}
```

## 4. Alternative Tools

- **Mutex (Mutual Exclusion):** Similar to a Critical Section, but it works across *different programs* (e.g., syncing Word and Excel). Slower than Critical Sections.
- **Semaphores:** A counter. Instead of "Only 1 person," it allows "Up to 5 people." Good for limiting database connections.
- **Events:** A starting gun. Thread A sleeps until Thread B fires the "Event" signal.

## 5. Summary Checklist

1. **Sleep** pauses a thread to save CPU.
2. **Race Conditions** corrupt data when threads fight over memory.
3. **Critical Sections** lock a piece of code so only one thread runs it at a time.
4. **Enter/Leave:** You must always pair EnterCriticalSection with LeaveCriticalSection. If you forget to Leave, the program hangs forever (Deadlock).

# UNDERSTANDING CRITICAL SECTIONS

Critical sections are synchronization mechanisms that enforce exclusive access to shared resources or code blocks by multiple threads within a process. This prevents race conditions and ensures data consistency.

## 1. Key Functions

Initializing a Critical Section:

```
#include <windows.h>

CRITICAL_SECTION cs; // Declare a global critical section object

// Initialize the critical section before use:
InitializeCriticalSection(&cs);
```

Entering a Critical Section:

```
EnterCriticalSection(&cs); // Acquire ownership of the critical section

// Code that accesses shared resources or executes critical code

LeaveCriticalSection(&cs); // Release ownership
```

Deleting a Critical Section:

```
// When no longer needed:
DeleteCriticalSection(&cs); // Free up associated resources
```



BigJob1  
program.mp4

## Events and Thread Signaling (BIGJOB1 & BIGJOB2)

We now understand how to *lock* threads (CriticalSection). Now, let's learn how to *talk* to them.


### 1. The Core Concept: Signaling

In a multithreaded app, threads often need to wait for each other.

- **The Main Thread:** "I need the Worker Thread to start calculating now."
- **The Worker Thread:** "I am finished. Here is the result."

We could use a while loop to constantly check a variable (Polling), but that wastes 100% of the CPU. Instead, we use **Event Objects**. An Event Object is like a **Start Gun**.

- **Unsignaled State:** The gun is raised. Everyone waits.
- **Signaled State:** *BANG!* The gun fires. The waiting thread wakes up and runs immediately.



## 2. The BIGJOB1 Program (The "Create/Destroy" Method)


This program calculates a complex math benchmark ("Savage").

### The Strategy:

1. **Start:** When you click "Start", the Main Thread calls `_beginthread`.
2. **Run:** The new thread runs the calculation loop.
3. **Finish:** When the calculation is done, the thread sends a message (`WM_CALC_DONE`) to the Main Window and then **destroys itself** (`_endthread`).

**The Communication Trick:** How does the Main Thread stop the calculation if the user clicks "Cancel"? It sets a boolean flag `bContinue = FALSE`. The Worker Thread checks this flag inside its for loop. If it sees FALSE, it aborts.

**The Flaw:** Creating and destroying a thread every time you click "Start" is slow (Thread Creation Overhead). It is better to keep the thread alive and just put it to sleep.



### 3. The BIGJOB2 Program (The "Event" Method)


This version is much smarter. It creates the Worker Thread *once* when the program starts.

#### The Strategy:

1. **Sleep:** The Worker Thread sits in an infinite loop, waiting at a WaitForSingleObject call. It consumes 0% CPU.
2. **Wake Up:** When you click "Start", the Main Thread calls SetEvent.
3. **Run:** The WaitForSingleObject sees the signal, wakes up, runs the calculation *once*, and then loops back to the top to wait again.

#### The Key Functions:

- **CreateEvent:** Makes the object.
  - ✓ *Auto-Reset:* Once one thread wakes up, the event automatically turns off (Unsignaled) so the thread doesn't run twice.
  - ✓ *Manual-Reset:* The event stays On until you manually turn it off.
- **SetEvent:** Fires the gun. Wakes up the waiting thread.
- **ResetEvent:** Reloads the gun (makes it Unsigned).
- **WaitForSingleObject:** The thread pauses here until the Event is Signaled.



## 4. Comparison: Polling vs. Events

Bad (Polling):

```
// Worker Thread
while (bNotDone) {
    if (bReady) { DoWork(); } // Wastes CPU checking this 1,000,000 times/sec
}
```

Good (Events):

```
// Worker Thread
while (TRUE) {
    WaitForSingleObject(hEvent, INFINITE); // Sleeps. CPU usage = 0%
    DoWork();
}
```

## 5. Summary Checklist

1. **Events** are the best way to coordinate threads without wasting CPU.
2. **SetEvent** wakes up a waiting thread.
3. **WaitForSingleObject** puts a thread to sleep until the signal comes.
4. **Auto-Reset Events** are safer for "Start Work" signals because they turn themselves off automatically.
5. **BIGJOB2** is better than BIGJOB1 because it reuses the thread instead of recreating it.

# PERSISTENT THREAD MODEL

## Why a Persistent Thread?

Instead of creating and destroying threads repeatedly, the program keeps **one worker thread alive for the entire lifetime** of the app.

### I. Why This Model Is Used

This approach avoids the cost of repeatedly creating and destroying threads and provides a more stable execution environment for background work.

### II. Practical Benefits

- **Reduced overhead:** Thread creation and teardown are expensive operations.
- **Improved responsiveness:** The thread is already available when work begins.
- **Predictable performance:** Memory and scheduling behavior remain consistent.

### III. Ideal Use Cases

- Repeated or long-running background tasks
- Work triggered by user input or window messages
- Applications where startup latency matters

## Synchronization Reality

### I. The Moment Things Get Dangerous

A persistent thread is simple **only** when it works alone. The moment shared data is introduced, synchronization becomes mandatory.

### II. Shared State Problems

When multiple threads access shared variables:

- Race conditions become unavoidable without protection.
- Bugs become timing-dependent and hard to reproduce.

### III. Required Synchronization Tools

- Critical sections
- Mutexes
- Events

## IV. Rule of Thumb

One thread → simple

Two threads touching shared data → synchronize or suffer

# EVENT-DRIVEN ARCHITECTURE

## 1. Core Design Principle

### I. Event-Driven vs Polling

The architecture is event-driven, meaning threads sleep until something meaningful happens, rather than constantly checking conditions.

### II. Signal Flow

- The **window procedure** signals an event when work should begin.
- The **worker thread** waits on that event.
- While waiting, the thread consumes no CPU time.
- Once signaled, the thread wakes up, performs its task, and returns to waiting.

### III. Benefits of This Model

- Zero busy-waiting
- Clean separation between UI and background work
- High responsiveness even under load

## 2. Event Types and Their Meaning

### I. Auto-Reset Events

- Automatically reset after waking one thread
- Ideal for one-time notifications
- Common in producer/consumer patterns

### II. Manual-Reset Events

- Remain signaled until explicitly reset
- Useful when multiple actions depend on the same signal
- Require careful control to avoid logic errors

### **III. Design Awareness**

Choosing the correct event type is not optional — it directly affects correctness and performance.

## **USER INTERACTION AND STATUS FEEDBACK**

### **1. Input Design**

#### **I. Mouse-Driven Control**

- **Left mouse click:** Starts the calculation
- **Right mouse click:** Aborts the calculation

This design is minimal, intentional, and avoids unnecessary UI complexity.

### **2. Status Communication**

#### **I. Information Presented to the User**

- Current program state (idle, running, aborted)
- Time taken for calculations

#### **II. Why This Matters**

A responsive UI is not a luxury. If users cannot tell what the program is doing, the program is effectively broken.

### **III. Possible Enhancements**

- Progress bars
- Activity indicators
- Periodic progress messages from the worker thread

# CODE STRUCTURE AND RESPONSIBILITIES 🧠

## 1. Window Procedure (WndProc)

### I. Primary Responsibilities

- Handle window creation and cleanup
- Process mouse input
- Signal events to the worker thread
- Receive completion or abort messages
- Paint status information

### II. Mental Model

The window procedure is a **coordinator**, not a worker.

## 2. Worker Thread Logic

### I. Core Execution Pattern

- Infinite loop
- Wait for event
- Perform calculation
- Check a control flag (bContinue) to allow abortion
- Notify the window of completion or cancellation

### II. Mental Model

The worker thread is a **sleeping worker**, not a CPU spinner.

## KEY TAKEAWAYS 🧠

### 1. Architectural Lessons

- Persistent threads improve responsiveness and efficiency
- Events are the backbone of clean WinAPI multithreading
- UI responsiveness is part of program correctness

### 2. Design Discipline

- UI logic and worker logic must remain separated
- Abort paths must be treated as first-class logic
- Sleeping threads are superior to polling threads

## THREAD LOCAL STORAGE (TLS) 💡

### 1. Purpose of TLS

#### I. What TLS Provides

Thread Local Storage allows each thread to store its own private data, even though all threads execute the same code.

#### II. When TLS Is Necessary

- Each thread needs its own state
- Global variables would conflict
- Passing context through function parameters becomes unmanageable

### 2. TLS Lifecycle

#### I. Define a Data Structure

Create a structure representing the per-thread data.

#### II. Allocate a TLS Index

Use TlsAlloc once and store the index globally.

#### III. Assign Data Per Thread

Allocate memory for the structure and associate it with the TLS index using TlsSetValue.

## **IV. Access TLS Data**

Retrieve the data anywhere in the thread using TlsGetValue.

## **V. Cleanup Per Thread**

Free the allocated memory when the thread exits.

## **VI. Free the TLS Index**

Release the TLS index after all threads are finished.

## **3. TLS Rules That Matter**

### **I. Isolation Advantage**

TLS data is thread-private, so locking is unnecessary for that data.

### **II. Responsibility Reminder**

TLS does not manage memory for you — allocation and cleanup are still your job.

### **III. Common Failure**

Forgetting cleanup leads to silent memory leaks that are difficult to trace.

## **FINAL MENTAL MODEL**

### **1. Core Ideas to Remember**

- **Persistent thread** → always ready
- **Event-driven signaling** → no wasted CPU
- **WndProc** → coordinator
- **Worker thread** → sleeper
- **TLS** → per-thread memory, not shared state

This is **real WinAPI architecture**, not textbook decoration.

# THREAD LOCAL STORAGE (TLS): ENABLING THREAD-SPECIFIC DATA IN MULTITHREADED ENVIRONMENTS

## 1. Understanding the Need for TLS

### I. Shared Memory Reality in Multithreaded Programs

In a multithreaded application, threads typically share:

- Global variables
- Heap-allocated memory

This shared access is powerful, but it also introduces risk.

### II. The Core Problem TLS Solves

Some data **must not be shared**, even though the code accessing it is identical across threads. Examples include:

- Per-thread counters
- Thread-specific state or context
- Temporary buffers used during calculations

If such data is shared, threads can overwrite each other's values, leading to corruption and unpredictable behavior.

### III. Why Locks Are Not Always the Answer

Using critical sections or mutexes for every access:

- Adds overhead
- Complicates logic
- Can still lead to deadlocks if misused

TLS avoids this entirely by giving **each thread its own private copy of data**.

### IV. What TLS Provides

Thread Local Storage allows each thread to associate data with itself, rather than with the process as a whole.

Each thread sees **only its own value**, even though all threads use the same TLS index.

## 2. Conceptual Model of TLS

### I. TLS as a Per-Thread Slot Table

You can think of TLS as:

- A table indexed by a TLS index
- Each thread has its own version of this table

The same index points to **different data** depending on which thread is running.

### II. Key Design Advantage

- No synchronization required for TLS data
- No accidental cross-thread modification
- Clean separation of thread-specific state

## 3. Windows API Functions for TLS Management

### I. TlsAlloc

Allocates a free TLS index.

- The returned index is process-wide
- It acts as an identifier for thread-local data

This is typically done once during program initialization.

### II. TlsSetValue

Associates a value with a TLS index **for the calling thread only**.

- The value is usually a pointer to a dynamically allocated structure
- Other threads using the same index will not see this value

This step is usually performed when a thread starts.

### **III. TlsGetValue**

Retrieves the value associated with a TLS index **for the calling thread**.

- Safe to call from anywhere in the thread
- Returns NULL if no value has been set

This allows thread-specific data to be accessed without passing parameters through function calls.

### **IV. TlsFree**

Releases a previously allocated TLS index.

- Should be called only after all threads are done using it
- Does not automatically free per-thread memory

Freeing the TLS index too early can cause undefined behavior.

## **4. Lifetime and Responsibility Considerations**

### **I. Memory Management Is Manual**

TLS does **not** manage memory for you.

- You allocate memory for each thread
- You must free it when the thread exits

Failure to do so results in memory leaks that are easy to miss.

### **II. Thread Exit Cleanup**

Each thread should:

- Retrieve its TLS value
- Free any allocated memory
- Optionally clear the TLS slot

This is commonly done just before the thread terminates.

## 5. When TLS Is the Right Tool

### I. Ideal Use Cases

- Per-thread error information
- Thread-specific caches
- Independent execution contexts
- Libraries that must be thread-safe without global locks

### II. When TLS Is Overkill

- Data that must be shared anyway
- Simple single-threaded applications
- Cases where passing parameters is simpler and clearer

## 6. Key Takeaways 🧠

### I. What TLS Gives You

- True thread isolation
- Cleaner multithreaded design
- Reduced synchronization complexity

### II. What TLS Does Not Do

- It does not replace synchronization for shared data
- It does not manage memory automatically
- It does not protect poorly designed thread logic

### III. Mental Rule

Shared data → synchronize

Thread-specific data → TLS

## 7. Example Code

(Implementation follows to demonstrate allocation, assignment, access, and cleanup of TLS data in a multithreaded WinAPI application.)

```
1 // Allocate a TLS index for thread-specific data
2 DWORD dwTlsIndex = TlsAlloc();
3
4 // Within each thread:
5 PDATA pdata = (PDATA)GlobalAlloc(GPTR, sizeof(DATA));
6 pdata->a = /* thread-specific value */;
7 pdata->b = /* another thread-specific value */;
8 TlsSetValue(dwTlsIndex, pdata);
9
10 // Access thread-local data:
11 PDATA pdata = (PDATA)TlsGetValue(dwTlsIndex);
12 // Use pdata->a and pdata->b as needed
13
14 // Before thread termination:
15 GlobalFree(TlsGetValue(dwTlsIndex));
16
17 // When all threads using the TLS index are done:
18 TlsFree(dwTlsIndex);
```

## THREAD LOCAL STORAGE (THE EASY WAY)

We previously learned that Global variables are dangerous (shared by everyone) and Local variables are temporary (die when the function ends).

**Thread Local Storage (TLS)** is the magic middle ground: A global variable that gives each thread its own private copy.

The old way (using TlsAlloc, TlsGetValue) is painful. Microsoft gave us a shortcut.

### 1. The Magic Keyword: `_declspec(thread)`

This is a compiler trick that handles all the hard work for you. You don't need to call API functions. You just declare the variable.

**The Syntax:**

```
_declspec(thread) int iGlobal = 1; // Lives globally, but unique per thread
```

**How it works:**

- **Thread A** reads iGlobal, it sees **1**. It adds 5. Now it sees **6**.
- **Thread B** reads iGlobal, it still sees **1**. It adds 100. Now it sees **101**.
- **Thread A** reads iGlobal again, it still sees **6**.

## 2. Initialization and Usage

The compiler automatically gives every new thread its own fresh copy of the variable when the thread starts.

### Example 1: Simple Integer

```
#include <windows.h>
#include <stdio.h>

// 1. Declare it with the magic keyword
__declspec(thread) int iCount = 0;

void ThreadFunc() {
    // 2. Use it like a normal variable
    iCount++;
    printf("Thread ID: %d, Count: %d\n", GetCurrentThreadId(), iCount);
}

// If you run 3 threads, they ALL print "Count: 1". None of them affect the others.
```

**Example 2: Complex Structs (Pointers)** You can use it for pointers too, but you must allocate the memory yourself inside the thread.

```
typedef struct { int id; char* name; } UserData;

// The POINTER is thread-local.
// Thread A has its own pointer, Thread B has its own pointer.
__declspec(thread) UserData* pData = NULL;

void ThreadFunc() {
    // Each thread allocates its own memory block
    pData = (UserData*)malloc(sizeof(UserData));
    pData->id = GetCurrentThreadId();
    // ... use it ...
}
```


### 3. The Catch (Limitations)

This magic comes with rules:


1. **Microsoft Only:** This is not standard C++. It won't work on Linux or GCC (unless they have similar extensions).
2. **No Complex Constructors:** In C++, you generally cannot use this for classes that need complex initialization code (constructors/destructors) to run automatically. It works best for simple types (int, float, pointers).
3. **DLL Issues:** If you are writing a DLL, `__declspec(thread)` can cause crashes on older Windows versions (pre-Vista) if the DLL is loaded dynamically (`LoadLibrary`).

### Summary Checklist for Chapter 20

1. **Multithreading** allows apps to do two things at once (UI + Math).
2. **\_beginthread** is better than `CreateThread` for C programs.
3. **Race Conditions** kill programs. Use **Critical Sections** to stop them.
4. **Events** allow threads to signal each other ("I'm done!").
5. **TLS (`__declspec(thread)`)** gives every thread its own private "global" variable.



# CHAPTER 20: THREAD LOCAL STORAGE



Petzold, Programming Windows (6<sup>th</sup> Edition)

